epoll原理简介

通过上面的分析,poll运行效率的两个瓶颈已经找出,现在的问题是怎么改进。首先,每次poll都要把 1000个fd 拷入内核,太不科学了,内核干嘛不自己保存已经拷入的fd呢?答对了,epoll就是自己保存拷入的fd,它的API就已经说明了这一点——不是 epoll_wait的时候才传入fd,而是通过epoll_ctl把所有fd 传入内核再一起"wait",这就省掉了不必要的重复拷贝。其次,在 epoll_wait时,也不是把current轮流的加入fd对应的设备等待队列,而是在设备等待队列醒来时调用一个回调函数(当然,这就需要"唤醒回调"机制),把产生事件的fd归入一个链表,然后返回这个链表上的fd。

epoll剖析

```
epoll是个module,所以先看看module的入口eventpoll init
[fs/eventpoll.c-->evetpoll_init()]
 1582 static int __init eventpoll_init(void)
 1583 {
 1584
         int error;
 1585
 1586
         init_MUTEX(&epsem);
 1587
 1588
         /* Initialize the structure used to perform safe poll wait head wake ups */
 1589
         ep poll safewake init(&psw);
 1590
 1591
         /* Allocates slab cache used to allocate "struct epitem" items */
 1592
         epi_cache = kmem_cache_create("eventpoll_epi", sizeof(struct epitem),
 1593
               0, SLAB_HWCACHE_ALIGN|EPI_SLAB_DEBUG|SLAB_PANIC,
 1594
              NULL, NULL);
 1595
         /* Allocates slab cache used to allocate "struct eppoll entry" */
 1596
 1597
         pwq_cache = kmem_cache_create("eventpoll_pwq",
 1598
              sizeof(struct eppoll_entry), 0,
 1599
              EPI_SLAB_DEBUG|SLAB_PANIC, NULL, NULL);
 1600
 1601
          * Register the virtual file system that will be the source of inodes
 1602
 1603
          * for the eventpoll files
 1604
         error = register_filesystem(&eventpoll_fs_type);
 1605
 1606
         if (error)
 1607
            goto epanic;
 1608
 1609
         /* Mount the above commented virtual file system */
 1610
         eventpoll_mnt = kern_mount(&eventpoll_fs_type);
 1611
         error = PTR ERR(eventpoll mnt);
         if (IS ERR(eventpoll mnt))
 1612
 1613
            goto epanic;
 1614
 1615
         DNPRINTK(3, (KERN_INFO "[%p] eventpoll: successfully initialized.\n",
 1616
              current));
 1617
         return 0;
 1618
```

```
1619 epanic:
       panic("eventpoll_init() failed\n");
 1620
 1621 }
很有趣,这个module在初始化时注册了一个新的文件系统,叫"eventpollfs"(在eventpoll_fs_type结
构里),然后挂载此文件系统。另外创建两个内核cache(在内核编程中,如果需要频繁分配小块内存
应该创建kmem_cahe来做"内存池"),分别用于存放struct epitem和eppoll_entry。如果以后要开发新
的文件系统,可以参考这段代码。
现在想想epoll create为什么会返回一个新的fd?因为它就是在这个叫做"eventpollfs"的文件系统里创建
了一个新文件!如下:
[fs/eventpoll.c-->sys_epoll_create()]
 476 asmlinkage long sys_epoll_create(int size)
 477 {
 478
       int error, fd;
 479
       struct inode *inode;
 480
       struct file *file;
 481
 482
       DNPRINTK(3, (KERN INFO "[%p] eventpoll: sys epoll create(%d)\n",
 483
            current, size));
 484
 485
       /* Sanity check on the size parameter */
 486
       error = -EINVAL;
 487
       if (size \leq 0)
 488
         goto eexit 1;
 489
 490
 491
        * Creates all the items needed to setup an eventpoll file. That is,
 492
        * a file structure, and inode and a free file descriptor.
 493
 494
       error = ep_getfd(&fd, &inode, &file);
 495
       if (error)
 496
         goto eexit_1;
 497
 498
       /* Setup the file internal data structure ( "struct eventpoll" ) */
 499
       error = ep file init(file);
 500
       if (error)
 501
         goto eexit_2;
函数很简单,其中ep_getfd看上去是"get",其实在第一次调用epoll_create时,它是要创建新inode、
新的file、新的fd。而ep file init则要创建一个struct eventpoll结构,并把它放入file-
>private data,注意,这个private data后面还要用到的。
看到这里,也许有人要问了,为什么epoll的开发者不做一个内核的超级大map把用户要创建的epoll句柄
存起来,在epoll_create时返回一个指针?那似乎很直观呀。但是,仔细看看,linux的系统调用有多少是
返回指针的?你会发现几乎没有!(特此强调,malloc不是系统调用,malloc调用的brk才是)因为linux
做为unix的最杰出的继承人,它遵循了unix的一个巨大优点——一切皆文件,输入输出是文件、socket也
是文件,一切皆文件意味着使用这个操作系统的程序可以非常简单,因为一切都是文件操作而已!(unix
还不是完全做到,plan 9才算)。而且使用文件系统有个好处:epoll_create返回的是一个fd,而不是该
死的指针,指针如果指错了,你简直没办法判断,而fd则可以通过current->files->fd array[]找到其真
epoll_create好了,该epoll_ctl了,我们略去判断性的代码:
[fs/eventpoll.c-->sys epoll ctl()]
 524 asmlinkage long
 525 sys_epoll_ctl(int epfd, int op, int fd, struct epoll_event __user *event)
 526 {
```

```
527
       int error;
 528
       struct file *file, *tfile;
 529
       struct eventpoll *ep;
 530
       struct epitem *epi;
 531
       struct epoll_event epds;
 575
       epi = ep find(ep, tfile, fd);
 576
       error = -EINVAL;
 577
 578
       switch (op) {
 579
       case EPOLL_CTL_ADD:
 580
          if (!epi) {
 581
            epds.events |= POLLERR | POLLHUP;
 582
 583
            error = ep_insert(ep, &epds, tfile, fd);
 584
          } else
 585
            error = -EEXIST;
 586
          break:
 587
       case EPOLL_CTL_DEL:
 588
          if (epi)
 589
            error = ep_remove(ep, epi);
 590
          else
 591
            error = -ENOENT;
 592
          break;
 593
       case EPOLL CTL MOD:
 594
          if (epi) {
 595
            epds.events |= POLLERR | POLLHUP;
 596
            error = ep_modify(ep, epi, &epds);
 597
 598
            error = -ENOENT;
 599
          break;
 600
原来就是在一个大的结构(现在先不管是什么大结构)里先ep_find,如果找到了struct epitem而用户操
作是ADD,那么返回-EEXIST;如果是DEL,则ep remove。如果找不到struct epitem而用户操作是
ADD,就ep_insert创建并插入一个。很直白。那这个"大结构"是什么呢?看ep_find的调用方式,ep参数
应该是指向这个"大结构"的指针,再看ep = file->private data,我们才明白,原来这个"大结构"就是那
个在epoll_create时创建的struct eventpoll,具体再看看ep_find的实现,发现原来是struct eventpoll
的rbr成员(struct rb_root),原来这是一个红黑树的根!而红黑树上挂的都是struct epitem。
现在清楚了,一个新创建的epoll文件带有一个struct eventpoll结构,这个结构上再挂一个红黑树,而这
个红黑树就是每次epoll ctl时fd存放的地方!
现在数据结构都已经清楚了,我们来看最核心的:
[fs/eventpoll.c-->sys_epoll_wait()]
 627 asmlinkage long sys_epoll_wait(int epfd, struct epoll_event __user *events,
 628
                int maxevents, int timeout)
 629 {
 630
       int error;
       struct file *file;
 631
 632
       struct eventpoll *ep;
 633
 634
       DNPRINTK(3, (KERN_INFO "[%p] eventpoll: sys_epoll_wait(%d, %p, %d, %d)\n",
 635
             current, epfd, events, maxevents, timeout));
 636
 637
       /* The maximum number of event must be greater than zero */
```

```
638
         if (maxevents \leq 0)
  639
            return -EINVAL;
  640
  641
         /* Verify that the area passed by the user is writeable */
         if ((error = verify_area(VERIFY_WRITE, events, maxevents * sizeof(struct
  642
epoll_event))))
  643
           goto eexit_1;
  644
  645
         /* Get the "struct file *" for the eventpoll file */
  646
         error = -EBADF;
  647
         file = fget(epfd);
  648
         if (!file)
  649
           goto eexit_1;
  650
  651
         * We have to check that the file structure underneath the fd
  652
  653
          * the user passed to us _is_ an eventpoll file.
  654
          */
  655
         error = -EINVAL;
  656
         if (!IS_FILE_EPOLL(file))
  657
           goto eexit_2;
  658
  659
          * At this point it is safe to assume that the "private_data" contains
  660
         * our own data structure.
  661
  662
  663
         ep = file->private_data;
  664
         /* Time to fish for events ... */
  665
  666
         error = ep_poll(ep, events, maxevents, timeout);
  667
  668 eexit_2:
  669
         fput(file);
  670 eexit 1:
         DNPRINTK(3, (KERN_INFO "[%p] eventpoll: sys_epoll_wait(%d, %p, %d, %d) =
  671
%d\n",
  672
               current, epfd, events, maxevents, timeout, error));
  673
  674
         return error;
  675 }
故伎重演,从file->private_data中拿到struct eventpoll,再调用ep_poll
[fs/eventpoll.c-->sys_epoll_wait()->ep_poll()]
 1468 static int ep_poll(struct eventpoll *ep, struct epoll_event __user *events,
 1469
              int maxevents, long timeout)
 1470 {
 1471
         int res, eavail;
 1472
         unsigned long flags;
 1473
         long jtimeout;
 1474
         wait_queue_t wait;
 1475
 1476
 1477
          * Calculate the timeout by checking for the "infinite" value ( -1 )
 1478
          * and the overflow condition. The passed timeout is in milliseconds,
```

```
1479
         * that why (t * HZ) / 1000.
 1480
         itimeout = timeout == -1 | timeout > (MAX SCHEDULE TIMEOUT - 1000) / HZ?
 1481
 1482
           MAX_SCHEDULE_TIMEOUT: (timeout * HZ + 999) / 1000;
 1483
 1484 retry:
 1485
         write lock irgsave(&ep->lock, flags);
 1486
         res = 0:
 1487
 1488
         if (list_empty(&ep->rdllist)) {
 1489
           /*
            * We don't have any available event to return to the caller.
 1490
            * We need to sleep here, and we will be wake up by
 1491
 1492
            * ep_poll_callback() when events will become available.
 1493
 1494
           init_waitqueue_entry(&wait, current);
 1495
           add wait queue(&ep->wg, &wait);
 1496
 1497
           for (;;) {
 1498
              /*
               * We don't want to sleep if the ep_poll_callback() sends us
 1499
 1500
               * a wakeup in between. That's why we set the task state
               * to TASK INTERRUPTIBLE before doing the checks.
 1501
 1502
 1503
              set current state(TASK INTERRUPTIBLE);
 1504
              if (!list_empty(&ep->rdllist) || !jtimeout)
 1505
                 break;
 1506
              if (signal_pending(current)) {
 1507
                res = -EINTR;
 1508
                break;
 1509
              }
 1510
 1511
              write_unlock_irgrestore(&ep->lock, flags);
 1512
              itimeout = schedule timeout(jtimeout);
 1513
              write_lock_irqsave(&ep->lock, flags);
 1514
 1515
           remove_wait_queue(&ep->wq, &wait);
 1516
 1517
           set current state(TASK RUNNING);
         }
 1518
又是一个大循环,不过这个大循环比poll的那个好,因为仔细一看——它居然除了睡觉和判断ep->rdllist
是否为空以外,啥也没做!
什么也没做当然效率高了,但到底是谁来让ep->rdllist不为空呢?
答案是ep insert时设下的回调函数:
[fs/eventpoll.c-->sys_epoll_ctl()-->ep_insert()]
 923 static int ep insert(struct eventpoll *ep, struct epoll event *event,
 924
              struct file *tfile, int fd)
 925 {
 926
        int error, revents, pwake = 0;
 927
        unsigned long flags;
 928
        struct epitem *epi;
 929
        struct ep_pqueue epq;
```

```
930
 931
        error = -ENOMEM;
 932
        if (!(epi = EPI_MEM_ALLOC()))
 933
           goto eexit_1;
 934
 935
        /* Item initialization follow here ... */
 936
        EP RB INITNODE(&epi->rbn);
 937
        INIT LIST HEAD(&epi->rdllink);
 938
        INIT_LIST_HEAD(&epi->fllink);
 939
        INIT_LIST_HEAD(&epi->txlink);
 940
        INIT_LIST_HEAD(&epi->pwqlist);
 941
        epi->ep = ep;
 942
        EP_SET_FFD(&epi->ffd, tfile, fd);
 943
        epi->event = *event;
 944
        atomic_set(&epi->usecnt, 1);
 945
        epi->nwait = 0;
 946
 947
        /* Initialize the poll table using the gueue callback */
 948
        epq.epi = epi;
 949
        init_poll_funcptr(&epq.pt, ep_ptable_queue_proc);
 950
 951
         * Attach the item to the poll hooks and get current event bits.
 952
 953
         * We can safely use the file* here because its usage count has
 954
         * been increased by the caller of this function.
 955
         */
 956
        revents = tfile->f_op->poll(tfile, &epq.pt);
我们注意949行,其实就是
&(epq.pt)->qproc = ep_ptable_queue_proc;
紧接着 tfile->f_op->poll(tfile, &epq.pt)其实就是调用被监控文件 ( epoll里叫"target file")的poll方
法,而这个poll其实就是调用poll_wait(还记得poll_wait吗?每个支持poll的设备驱动程序都要调用
的),最后就是调用ep_ptable_queue_proc。这是比较难解的一个调用关系,因为不是语言级的直接调
ep insert还把struct epitem放到struct file里的f ep links连表里,以方便查找,struct epitem里的
fllink就是担负这个使命的。
[fs/eventpoll.c-->ep ptable queue proc()]
 883 static void ep_ptable_queue_proc(struct file *file, wait_queue_head_t *whead,
 884
                poll_table *pt)
 885 {
 886
        struct epitem *epi = EP ITEM FROM EPOUEUE(pt);
 887
        struct eppoll_entry *pwq;
 888
 889
        if (epi->nwait >= 0 \&\& (pwq = PWQ_MEM_ALLOC())) {
 890
           init_waitqueue_func_entry(&pwq->wait, ep_poll_callback);
 891
           pwg->whead = whead;
 892
           pwq->base = epi;
 893
           add wait queue(whead, &pwg->wait);
 894
           list_add_tail(&pwq->llink, &epi->pwqlist);
 895
           epi->nwait++;
 896
        } else {
 897
           /* We have to signal that an error occurred */
 898
           epi->nwait = -1;
 899
```

```
900 }
上面的代码就是ep_insert中要做的最重要的事:创建struct eppoll_entry,设置其唤醒回调函数为
ep poll callback, 然后加入设备等待队列(注意这里的whead就是上一章所说的每个设备驱动都要带的
等待队列)。只有这样,当设备就绪,唤醒等待队列上的等待着时,ep_poll_callback就会被调用。<mark>每次</mark>
调用poll系统调用,操作系统都要把current(当前进程)挂到fd对应的所有设备的等待队列上,可以想
象,fd多到上千的时候,这样"挂"法很费事;而每次调用epoll_wait则没有这么罗嗦,epoll只在epoll ctl
时把current挂一遍(这第一遍是免不了的)并给每个fd一个命令"好了就调回调函数",如果设备有事件
了,通过回调函数,会把fd放入rdllist,而每次调用epoll wait就只是收集rdllist里的fd就可以了
——epoll巧妙的利用回调函数,实现了更高效的事件驱动模型。
现在我们猜也能猜出来ep_poll_callback会干什么了——肯定是把红黑树上的收到event的epitem(代表
每个fd)插入ep->rdllist中,这样,当epoll_wait返回时,rdllist里就都是就绪的fd了!
[fs/eventpoll.c-->ep poll callback()]
 1206 static int ep_poll_callback(wait_queue_t *wait, unsigned mode, int sync, void *key)
 1207 {
 1208
        int pwake = 0;
 1209
        unsigned long flags;
 1210
        struct epitem *epi = EP ITEM FROM WAIT(wait);
 1211
        struct eventpoll *ep = epi->ep;
 1212
 1213
        DNPRINTK(3, (KERN_INFO "[%p] eventpoll: poll_callback(%p) epi=%p
ep=%p\n",
 1214
             current, epi->file, epi, ep));
 1215
        write_lock_irqsave(&ep->lock, flags);
 1216
 1217
 1218
         * If the event mask does not contain any poll(2) event, we consider the
 1219
 1220
         * descriptor to be disabled. This condition is likely the effect of the
 1221
         * EPOLLONESHOT bit that disables the descriptor when an event is received,
 1222
         * until the next EPOLL_CTL_MOD will be issued.
 1223
 1224
        if (!(epi->event.events & ~EP_PRIVATE_BITS))
 1225
          goto is_disabled;
 1226
        /* If this file is already in the ready list we exit soon */
 1227
 1228
        if (EP IS LINKED(&epi->rdllink))
 1229
          goto is_linked;
 1230
 1231
        list add tail(&epi->rdllink, &ep->rdllist);
 1232
 1233 is linked:
 1234
         * Wake up ( if active ) both the eventpoll wait list and the ->poll()
 1235
 1236
         * wait list.
 1237
 1238
        if (waitqueue_active(&ep->wq))
 1239
          wake up(\&ep->wq);
 1240
        if (waitqueue_active(&ep->poll_wait))
 1241
          pwake++;
 1242
 1243 is disabled:
        write_unlock_irgrestore(&ep->lock, flags);
 1244
 1245
```

```
/* We have to call this outside the lock */
1247 if (pwake)
1248 ep_poll_safewake(&psw, &ep->poll_wait);
1249
1250 return 1;
1251 }
```

真正重要的只有1231行的只一句,就是把struct epitem放到struct eventpoll的rdllist中去。现在我们可以画出epoll的核心数据结构图了:

